

Score written responses against a rubric in five minutes

A walkthrough of your first evaluation: install a local AI model, start the app, load a rubric, paste learner responses, run the evaluation and check how closely the model grades like you.

The scenario used throughout. You run a leadership programme. Four learners (Alex, Priya, Sam and Jordan) wrote a reflective diary about a situation at work that did not go as planned. You'll score each entry against a five-part rubric, give feedback, and check the model against your own judgement.

1 · Install

Ollama and the nimble model

2 · Start

python3 serve.py

3 · Connect

Settings → Test connection

4 · Rubric

Use the sample, or import from Excel

5 · Evaluate

Paste responses → Evaluate

6 · Review

Agree, change, or grade blind

What you need

- ☐ A Mac, Windows or Linux computer. For the recommended nimble model: about 18 GB of free memory while running (Apple Silicon with 32 GB+, or an NVIDIA card with 16 to 24 GB). Smaller machines can use `tev1` or a cloud model.
- ☐ Python 3.8 or newer (already on Mac and Linux; on Windows from python.org or the Microsoft Store).
- ☐ The app folder **nimble-rubric-evaluator** (unzipped).
- ☐ About 10 GB of disk space for the model download.

No model or install yet? You can still follow along: start the app (step 2), tick **Demo mode** in Settings, and skip to step 4. Demo mode invents plausible scores for practice.

1 Install Ollama and the nimble model

2 minutes plus download

Ollama runs AI models on your own computer. nimble is a "decision model": instead of writing text, it picks an answer and says how sure it is. That makes it good at choosing a rubric level.

1. Go to **ollama.com/download**, download Ollama for your computer and install it. You need **version 0.35 or newer**.
2. Open a terminal. Mac: Applications → Utilities → **Terminal**. Windows: Start → **PowerShell**.
3. Download the model (about 9.5 GB, once):

```
ollama pull nimble
```

Smaller computer? Use `ollama pull tev1` (4B) and choose "Ollama · tev1" in Settings later.

4. Check it worked:

```
ollama list
```

You should see `nimble:latest`. Ollama keeps running in the background. If it isn't, run `ollama serve`.

2 Start the app

30 seconds

1. Unzip **ollama-eval.zip**, for example onto your Desktop. You get a folder called **nimble-rubric-evaluator**.
2. In the terminal, go into the folder and start the app:

```
cd ~/Desktop/nimble-rubric-evaluator  
python3 serve.py
```

Windows: `cd $HOME\Desktop\nimble-rubric-evaluator` then `py serve.py`.

3. Your browser opens **http://localhost:8787**. If it doesn't, type that address in.
4. Keep the terminal window open while you work. Press **Ctrl+C** in it to stop the app.

Why serve.py? It's a tiny local web server (Python standard library, nothing to install). It shows the app and passes requests to Ollama, or to a cloud provider if you choose one, so your browser doesn't block them. Nothing is uploaded anywhere when you use local Ollama.

3 Connect to the model

15 seconds

Click ⚙️ **Settings** at the top.

The screenshot shows the 'Settings' page of the 'Nimble Rubric Evaluator'. The 'Connection' tab is active. At the top right, a badge shows 'nimble ready · Ollama 0.35.1'. The 'Connection' section has a dropdown menu set to 'Ollama · nimble (recommended)'. Below this, there are four input fields: 'Ollama URL' (set to '/ollama'), 'Model (scores)' (set to 'nimble'), 'Feedback model (optional, writes narrative feedback)' (set to 'e.g. qwen3:8b, llama3.1:8b, gemma3'), and 'keep_alive (Ollama)' (set to '10m'). A 'Test connection' button is highlighted. Below the fields, there is a checkbox for 'Demo mode (simulated scores, no model needed)' which is unchecked. At the bottom, a status bar shows 'Connected to Ollama 0.35.1. 2 model(s): nimble:latest, qwen3:8b.'.


- 1 **Connect to:** leave it on **Ollama · nimble (recommended)**.
- 2 **Ollama URL:** leave it as `/ollama` ("through serve.py to Ollama on this computer").
- 3 **Model:** `nimble`.
- 4 Click **Test connection**. 5 You should see "Connected" and nimble in the list.
- 6 The badge at the top right shows the connection on every page. A blue dot means ready.

If it says "Ollama offline": make sure the Ollama app is running (or run `ollama serve`) and that you opened the app through `python3 serve.py`. **If it says "nimble not installed":** run `ollama pull nimble`. **"No /v1/systemone endpoint":** update Ollama.

4 Look at the rubric

1 minute

Click  **Rubrics**. The sample rubric is already in your library. A rubric is a grid: **criteria** down the side, **levels** across the top, and a description of what each level looks like in every cell.



The screenshot shows the 'Edit' page for a rubric titled 'Reflective diary: workplace situation'. The interface includes a sidebar with a library of rubrics, a main editing area with a title, assignment description, and a grid of criteria and levels. The grid has four levels: Beginning (0 points), Developing (1 point), Proficient (2 points), and Exemplary (3 points). The criteria include 'Description of the situation', 'Feelings and self-awareness', 'Analysis and learning', 'Action plan', and 'Professional tone'. The 'Professional tone' criterion has a 'YES/NO CHECKS' section with 'NO LOOKS LIKE' and 'YES LOOKS LIKE' columns.

Library +  7

Edit Import & template

Reflective diary: workplace situation 4 criteria · 1 check · 4 levels

Export CSV Duplicate Delete Save

Assignment / question the learner answered (sent to Nimble with each response)

Write a reflective diary entry (200-400 words) about a recent situation at work that did not go as planned. Describe what happened, how you felt, what you learned, and what you will do differently next time.

Criterion	Beginning points 0 x	Developing points 1 x	Proficient points 2 x	Exemplary points 3 x
Description of the situation How clearly and objectively does the learner describe what happened? weight 1	The situation is missing, vague or impossible to follow.	The situation is described but key facts (who, what, when, why it mattered) are missing.	A clear account of what happened, who was involved and why it mattered.	A concise, well-structured and objective account with the relevant context and stakes.
Feelings and self-awareness How well does the learner examine their own feelings, assumptions and behaviour during the situation? weight 1	No mention of own feelings or role; blames others or circumstances only.	Names feelings but does not examine them or their effect on behaviour.	Explains feelings and how they influenced own actions, with some honesty about own role.	Candidly analyses own feelings, assumptions and behaviour, and their impact on others.
Analysis and learning How deeply does the learner analyse why things happened and draw transferable learning? weight 2	No analysis; restates events without explaining causes or lessons.	Some explanation of causes, but learning is generic (e.g. 'communication is important').	Identifies specific causes and states clear, relevant lessons learned.	Insightful analysis of causes from several perspectives, linked to principles or models, with transferable lessons.
Action plan How specific and realistic is the learner's plan for what they will do differently? weight 1	No plan or next step is given.	A vague intention such as 'I will try to do better'.	A specific action the learner will take in a similar situation.	Specific, realistic actions with when and how they will be applied and how success will be judged.
Professional tone Is the entry written in a professional, respectful tone (no insults or confidential personal details about colleagues)? weight 0.5	NO LOOKS LIKE Contains insults, blame-laden language or confidential personal details about others.		YES LOOKS LIKE Respectful and professional throughout.	

+ Criterion + Level column + Yes/No check

4 criteria · 1 check · 4 levels · total weight 5.5 · rubric -839 of 8,192 tokens

Rubrics → Edit, showing the sample rubric

- 1 **Library**: click **Reflective diary: workplace situation**.
- 2 **Name** and 3 **Assignment**: the question the learner answered. It's sent with every response.
- 4 **Levels**, lowest to highest, each with points. 5 Each **row is a criterion** with a weight and a question for the model.
- 6 **Yes/no checks** underneath.
- 7 **Import & template** brings in your own rubric from Excel (see page 11).

The sample rubric in full


This is what the model is asked to judge. Notice each cell describes something you could point to in the writing, not just "good" or "poor". Analysis counts double (weight 2); tone counts half.

Criterion	Beginning (0)	Developing (1)	Proficient (2)	Exemplary (3)
Description of the situation How clearly and objectively does the learner describe what happened? weight 1	The situation is missing, vague or impossible to follow.	The situation is described but key facts (who, what, when, why it mattered) are missing.	A clear account of what happened, who was involved and why it mattered.	A concise, well-structured and objective account with the relevant context and stakes.
Feelings and self-awareness How well does the learner examine their own feelings, assumptions and behaviour during the situation? weight 1	No mention of own feelings or role; blames others or circumstances only.	Names feelings but does not examine them or their effect on behaviour.	Explains feelings and how they influenced own actions, with some honesty about own role.	Candidly analyses own feelings, assumptions and behaviour, and their impact on others.
Analysis and learning How deeply does the learner analyse why things happened and draw transferable learning? weight 2	No analysis; restates events without explaining causes or lessons.	Some explanation of causes, but learning is generic (e.g. 'communication is important').	Identifies specific causes and states clear, relevant lessons learned.	Insightful analysis of causes from several perspectives, linked to principles or models, with transferable lessons.
Action plan How specific and realistic is the learner's plan for what they will do differently? weight 1	No plan or next step is given.	A vague intention such as 'I will try to do better'.	A specific action the learner will take in a similar situation.	Specific, realistic actions with when and how they will be applied and how success will be judged.
Professional tone Is the entry written in a professional, respectful tone (no insults or confidential personal details about colleagues)? yes/no check · weight 0.5	NO: Contains insults, blame-laden language or confidential personal details about others.		YES: Respectful and professional throughout.	

5 Paste responses and run

1 minute

Click  **Evaluate**.



1 Reflective diary: workplace situation 4 criteria · 1 check · 4 levels View grid Edit in Rubrics

Learner responses Import CSV / TXT... Responses temp **2** Try sample Clear

Separate responses with a line containing only ---. Start a response with Name: Alex to label it, or import a CSV with Name and Response columns.

3 Name: Alex
Last Tuesday I was responsible for sending the quarterly report to our client, Brightwell, by 5pm. At 3pm I discovered that the sales figures from the regional team were missing, and I spent two hours chasing them instead of telling my manager. The report went out at 7pm and the client complained to our account director.

I felt panicked and embarrassed. Looking back, I think I avoided escalating because I didn't want to look like I couldn't cope, and that pride made the delay worse. My manager told me afterwards she could have called the regional lead directly and fixed it in ten minutes.

The real cause was not the missing data but that I had no checkpoint: I only checked inputs on the deadline day. I also assumed asking for help was a weakness, when in our team escalation is expected. This links to what we covered on "managing up" - surfacing risks early is part of the job.

Next quarter I will set a data-collection deadline three days before the report is due, put it in the shared calendar and if anything is missing 24 hours before I will message my manager straight away. I'll judge

4 4 responses

Name	Response	Words
Alex	Last Tuesday I was responsible for sending the quarterly report to our client, Brightwell, by 5...	211 words
Priya	I had a meeting with a customer and it did not go well. They were unhappy. I think communica...	37 words
Sam	The project kickoff was a disaster because Dave from IT is useless and never reads his email...	40 words
Jordan	During our team retrospective I was asked to facilitate for the first time. Halfway through, two...	132 words

Run

Assignment sent with each response (from the rubric)

Write a reflective diary entry (200-400 words) about a recent situation at work that did not go as planned. Describe what happened, how you felt, what you learned, and what you will do differently next time.

Run name (optional)

5 ☐ Narrative feedback (set a model in Settings)

☒ On-topic check

6 ☐ Blind grading

1 response is under 40 words and will be flagged for review.

7 Evaluate responses

8 Preview request

Evaluate, with the four sample diary entries loaded

- 1 Rubric:** choose **Reflective diary: workplace situation**.
- 2** Click **Try sample**.
- 3** The responses box fills with four diary entries.
- 4** Below it, the app shows how it split them: Alex, Priya, Sam and Jordan, with word counts.
- 5 Narrative feedback** needs a second, writing model (leave off). **On-topic check:** leave on. **6 Blind grading:** leave off for now (page 9).
- 7** Click **Evaluate responses**.
- 8 Preview request** shows exactly what is sent to the model.

Pasting your own responses

Separate responses with a line containing only ---. Start each with **Name: ...**. Or click **Import CSV / TXT** with Name and Response columns.

Name: Priya
I had a meeting with a customer and it did not go well.
They were unhappy. I think communication is important
and I will try to communicate better next time. Overall it
was a good learning experience.


Name: ...

What the model sees: the assignment text, then the learner's response, then one question per criterion with the level descriptions. It answers all of them together, in about a second per response on a fast computer. The first request can take up to a minute while the model loads.

6 Read the results

1 minute


When the run finishes, the app opens **Results**.



1 Summary: responses, average score, how many need a person to look at them, speed, and how many you've reviewed.

2 Charts: average per criterion and how many learners fell into each outcome band.

3 Blind grading switch. **4 Downloads:** CSV (Excel), full JSON, audit report, HTML report, print.



1 The model's level for each criterion. **Click the level name** to choose a different level if you disagree.

2 Click the circle to agree. It shows ✓ (agreed) or ✎ (you changed it).

3 Agree rest accepts all untouched scores on the row. **4** Tick rows, then use **5** the **bulk bar** to agree, set a level or clear for all of them. **6** Click a heading to sort.

What to expect from the sample: Alex's detailed entry should score high. Priya's three generic sentences should score low, and she is flagged as short. Sam blames a colleague by name, so "Professional tone" should be No. Every change shows a message in the bottom-left corner with **Undo** for 10 seconds (or Ctrl+Z).

Look closer at one learner

Click a learner's name, for example **Priya**.

Priya 67% Meets expectations needs review 0/5 reviewed

Possibly off-topic - Short (37 words)

Response 37 words - on-topic p:28% - 5 ms

I had a meeting with a customer and it did not go well. They were unhappy. I think communication is important and I will try to communicate better next time. Overall it was a good learning experience.

Scores Bars show the probability Nimble gave each level. Click a level to change it.

Description of the situation weight 1 - expected 1.86 of 3 - confidence 91%
A clear account of what happened, who was involved and why it mattered.

Feelings and self-awareness weight 1 - expected 2.61 of 3 - confidence 65%
Candidly analyses own feelings, assumptions and behaviour, and their impact on others.

Analysis and learning weight 2 - expected 2.23 of 3 - confidence 63%
Insightful analysis of causes from several perspectives, linked to principles or models, with transferable lessons.

Action plan weight 1 - expected 0.47 of 3 - confidence 81%
No plan or next step is given.

Professional tone weight 0.5 - confidence 52%
Contains insults, blame-laden language or confidential personal details about others.

Feedback for the learner built from the rubric level descriptions

What you did well:

- Description of the situation (Proficient): A clear account of what happened, who was involved and why it mattered.
- Feelings and self-awareness (Exemplary): Candidly analyses own feelings, assumptions and behaviour, and their impact on others.
- Analysis and learning (Exemplary): Insightful analysis of causes from several perspectives, linked to principles or models, with transferable lessons.

To improve:

- Description of the situation: you are at "Proficient". To reach "Exemplary": A concise, well-structured and objective account with the relevant context and stakes.
- Action plan: you are at "Beginning". To reach "Developing": A vague intention such as "I will try to do better".
- Professional tone: Respectful and professional throughout.

Audit trail

Response ID: w5uzkvgs

1 One button per level. The bar shows how likely the model thought each level was; the outlined button is the current grade. Click a level to change it.

2 **Agree with the rest** accepts every score you haven't reviewed.

3 **Feedback for the learner**, built from your rubric: what they did well, and exactly what the next level needs. Click **Copy feedback**.

4 **Audit trail**: the exact request, the model's raw reply with a SHA-256 fingerprint, the scoring arithmetic and every human change. Download it as HTML or JSON.

Use the < > buttons or the arrow keys to move between learners.

Why some results are flagged

- **Low confidence**: the model spread its probability across levels.
- **Borderline**: the top two levels were nearly tied.
- **Short**: under 40 words.
- **Possibly off-topic**: the on-topic check failed.

Flagged results are the ones a person should look at. Change the thresholds in Settings.

7 Check the model grades like you: blind grading

Before you trust the model with real learners, grade some responses yourself **without seeing its scores**, then compare. Turn on **Blind grading** on the Evaluate tab before a run, or with the switch at the top of Results.

Reflective diary: workplace situation · Oct 3, 2026, 9:49 PM
10/3/2026, 9:49:03 PM · rubric Reflective diary: workplace situation · model n1mb1e

☒ Blind grading

1 **Blind grading is on.** The model's scores, flags and feedback are hidden so they can't influence you. Grade each response by clicking **Grade ▼** in the table (or open a learner). When you're done, reveal the model's results to see how closely you agree.

4
ready to grade

0/4
responses with every criterion graded

6/20
30% done

Hidden
revealed when you click Reveal

Reveal model results

4

Responses

☐ Learner	Description of the situation	Feelings and self-awareness	Analysis and learning	Action plan	Professional tone	Graded	Words
☐ Alex	2 Proficient ▼ your grade	1 Developing ▼ your grade	Grade ▼ not graded	Grade ▼ not graded	Grade ▼ not graded	2/5	211
☐ Priya	3 Exemplary ▼ your grade	2 Proficient ▼ your grade	1 Developing ▼ your grade	Grade ▼ not graded	Grade ▼ not graded	3 3/5	37
☐ Sam	1 Developing ▼ your grade	Grade ▼ not graded	2 Grade ▼ not graded	Grade ▼ not graded	Grade ▼ not graded	1/5	40
☐ Jordan	Grade ▼ not graded	Grade ▼ not graded	Grade ▼ not graded	Grade ▼ not graded	Grade ▼ not graded	0/5	132

Tick the circle on a score to agree with the model, click the level name to pick a different one, or select rows to review in bulk. Click a learner for the full breakdown and audit trail. Every change can be undone for 10

Results with blind grading on

- 1** The dark banner means the model's scores, flags, charts, feedback and downloads are hidden.
- 2** Click **Grade ▼** and choose a level. **3** The **Graded** column tracks your progress.
- 4** When you've finished, click **Reveal model results**.

Read the alignment report

After you reveal, the **Model vs human alignment** report shows how closely the model matched you.

Model vs human alignment

How closely the model's scores match your grades. Updates as you review. [What do these numbers mean?](#)

8

All reviews

Blind grades only

7

Alignment report

1

Weak alignment

The model often disagrees with human grading. Improve the rubric wording (see Help) or keep a person grading.

2

SCORES COMPARED

20

4 responses - 20 graded blind

EXACT MATCH

15%

3 of 20 identical

WITHIN ONE LEVEL

44%

graded criteria only

MODEL HIGHER / LOWER

7 / 10

where you disagreed

OVERALL SCORE GAP

-13.6 pts

model minus final, per response

3

By criterion

Criterion	Compared	Exact match	Within 1 level	Model higher	Model lower	Average gap (levels)	Kappa	
Description of the situation	4	0%	50%	2	2	-0.25 model lower	needs 5+	Showing grid
Feelings and self-awareness	4	25%	75%	2	1	even	needs 5+	Show grid
Analysis and learning	4	0%	0%	2	2	even	needs 5+	Show grid
Action plan	4	0%	50%	0	4	-2.00 model lower	needs 5+	Show grid
Professional tone (yes/no)	4	50%	—	1	1	even	needs 5+	Show grid

4

Where you and the model landed: Description of the situation

Each number counts responses. The diagonal (outlined) is where you agreed; cells to the right of it mean the model scored higher than you.

You \ Model	Beginning	Developing	Proficient	Exemplary
Beginning				
Developing				1
Proficient				1
Exemplary	1		1	

5

By response

Learner	Reviewed	Matched model	You changed	Model overall	Final overall	Gap	
Alex blind	5/5	0	5	39%	48%	-9.1	Open ↓
Priya blind	5/5	0	5	67%	70%	-3.0	Open ↓
Sam blind	5/5	1	4	61%	30%	+30.3	Open ↓
Jordan blind	5/5	2	3	27%	100%	-72.7	Open ↓

6

Change history

22 changes, newest first. Every grade you give, change, clear or undo is recorded.

Results → Model vs human alignment

- 1 A plain-words **verdict**.
- 2 **Headline numbers**: exact match, within one level, whether the model scored higher or lower, and the overall gap.
- 3 **By criterion** and 4 the **grid**: your level down the side, the model's across the top; the outlined diagonal is agreement.
- 5 **By response**, with a per-learner download.
- 6 **Change history**: every grade you gave or changed, with the time and whether it was blind.
- 7 **Alignment report** downloads the whole thing to share.
- 8 Switch to **Blind grades only** for the fairest comparison.

Rule of thumb: grade at least 30 responses per rubric. Aim for 70%+ exact match on 4-level criteria, 90%+ within one level and 85%+ on yes/no checks. If one criterion is weak, rewrite its level descriptions to say what you'd see in the writing, and test again.

Nimble Rubric Evaluator · Quick start guide · page 10 of 13

8 Use your own rubric

Library
+ New

Reflective diary: workplace situation
4 criteria · 1 check · 4 levels

2

Drop a CSV file here
Or choose a file, or click here and paste (Ctrl+V / ⌘V) straight from Excel, Google Sheets or Word.

Choose file...
Paste as text instead

1

Start from a template
Same grid as the editor. The first rows are instructions and are ignored on import.

Blank template (CSV)
Filled sample

How to fill it in

3
'communication is important').","Identifies specific causes and states clear, relevant lessons learned.", "Insightful analysis of causes from several perspectives, linked to principles or models, with transferable lessons."
Action plan, score, 1, How specific and realistic is the learner's plan for what they will do differently?, No plan or next step is given., A vague intention such as 'I will try to do better', A specific action the learner will take in a similar situation., "Specific, realistic actions with when and how they will be applied and how success will be judged."
Professional tone, yes/no, 0.5, "Is the entry written in a professional, respectful tone (no insults or confidential personal details about colleagues)?", "Contains insults, blame-laden language or confidential personal details about others.", "Respectful and professional throughout."

Preview import

4
Preview: pasted rubric
5 ready

Name
Reflective diary: workplace situation

Assignment: Write a reflective diary entry (200-400 words) about a recent situation at work that did not go as planned. Describe what happened, how you felt, what you learned, and what you will do differently next time.

Row	Criterion	Beginning · 0	Developing · 1	Proficient · 2	Exemplary · 3
11	Description of the situation How clearly and objectively does the learner describe what happened? weight 1	The situation is missing, vague or impossible to follow.	The situation is described but key facts (who, what, when, why it mattered) are missing.	A clear account of what happened, who was involved and why it mattered.	A concise, well-structured and objective account with the relevant context and stakes.
12	Feelings and self-awareness How well does the learner examine their own feelings, assumptions and behaviour during the situation? weight 1	No mention of own feelings or role; blames others or circumstances only.	Names feelings but does not examine them or their effect on behaviour.	Explains feelings and how they influenced own actions, with some honesty about own role.	Candidly analyses own feelings, assumptions and behaviour, and their impact on others.
13	Analysis and learning How deeply does the learner analyse why things happened and draw transferable learning? weight 2	No analysis; restates events without explaining causes or lessons.	Some explanation of causes, but learning is generic (e.g. 'communication is important').	Identifies specific causes and states clear, relevant lessons learned.	Insightful analysis of causes from several perspectives, linked to principles or models, with transferable lessons.
14	Action plan How specific and realistic is the learner's plan for what they will do differently? weight 1	No plan or next step is given.	A vague intention such as 'I will try to do better'.	A specific action the learner will take in a similar situation.	Specific, realistic actions with when and how they will be applied and how success will be judged.
YES / NO CHECKS					
15	Professional tone Is the entry written in a professional, respectful tone (no insults or confidential personal details about colleagues)? weight 0.5	NO LOOKS LIKE Contains insults, blame-laden language or confidential personal details about others.	YES LOOKS LIKE Respectful and professional throughout.		

Cancel
Open in editor
Save to library

Rubrics → Import & template

- Download the **Blank template (CSV)** or **Filled sample** and open it in Excel or Google Sheets. Fill in one row per criterion: name, type (**score** or **yes/no**), weight, question for the model, then one column per level, lowest first, with headers like **Developing (1)**.
- Drop the saved CSV here, or **3** copy the cells and paste them.
- Check the preview, then click **Save to library**.

Other ways to connect a model

Under **Settings** → **Connect to**. Cloud options send learner text to that company: check your data policy first.

Option	Setup	Privacy	Probabilities
Ollama · nimble (recommended)	<code>ollama pull nimble</code>	Stays on your computer	Real
Ollama · tev1 / tev1:0.8b	<code>ollama pull tev1</code>	Local	Real
Ollama on another computer	Enter its address, e.g. <code>http://192.168.1.20:11434</code>	Your network	Real
TypeSafe Jev (cloud)	Paste your TypeSafe API key; model <code>jev-latest</code>	Sent to TypeSafe	Real
OpenAI, Anthropic, Google Gemini, OpenRouter, Groq, Mistral, other OpenAI-compatible	Paste the API key, Test connection, pick a model	Sent to provider	Model's own estimate
Ollama chat model, LM Studio	e.g. <code>ollama pull qwen3:8b</code>	Local	Model's own estimate

Connection

Choose which model scores your responses. Local Ollama + nimble is private and free. Cloud options send learner text to the provider. [How to choose and set up a model](#)

Connect to

1 TypeSafe Jev (cloud)

2 Jev decision API in the cloud. Same question format as nimble and real probabilities, but learner text is sent to the provider. Needs an API key. Requests go through `serve.py`.

API base URL

Model (scores)

3 API key (stored only in this browser)

Feedback model (optional, writes narrative feedback)

https://api.typesafe.ai/v1

jev-latest

paste your key Show

e.g. qwen3:8b, llama3.1:8b, gemma3

4 Test connection

☐ Demo mode (simulated scores, no model needed)

Settings with TypeSafe Jev chosen

- 1 Choose the option. 2 Read what it does with your data. 3 Paste the API key (stored only in this browser, never in exports). 4 Test connection.

A general LLM used as a judge gets the rubric as a prompt and replies in JSON; the app converts the reply so every screen works the same. Its "probabilities" are its own estimate, so review flags are rougher than nimble's.

Cheat sheet

Commands

```
# get the model, and see what's installed
ollama pull nimble
ollama list
# start Ollama if it isn't running
ollama serve
# start the app (add --port 9000 to change port)
python3 serve.py
# use Ollama on another machine
OLLAMA_HOST=http://server:11434 python3 serve.py
# allow another cloud API host
EXTRA_HOSTS=api.example.com python3 serve.py
```

Shortcuts

- **Ctrl+Z / ⌘Z**: undo the last action
- **← →** in a learner's detail: previous or next learner
- **Esc**: close a dialog or picker

Where things are kept

Rubrics and runs are saved in this browser on this computer. Use **History** → **Export all** to back up or move them. API keys stay in this browser only.

The full guide, with every screen explained, is in the app under **? Help**.

Troubleshooting

Ollama offline	Start Ollama; open the app via <code>python3 serve.py</code>
Model not installed	<code>ollama pull nimble</code>
No /v1/systemone	Update Ollama to 0.35+
401 / check API key	Paste the key again in Settings
429 rate limited	Settings → Requests in parallel = 1
Request too long	Shorten level text; split long responses
Scores too harsh	Add your graders' expectations to the level descriptions; check "model lower" in the alignment report
Deleted by mistake	Click Undo within 10 seconds

Writing rubrics that score well

- Describe what you'd see, not how good it is.
- One idea per criterion; 3 or 4 levels.
- Turn fuzzy judgements into yes/no checks.
- Write your graders' unwritten rules into the text.
- Test with blind grading and fix the weakest criterion.